

Unit-9. Concurrency Control Techniques

- Ankit Pangeni

DATE

Concurrency control is the procedure in DBMS for managing simultaneous operations without conflicting with each other. It is used to address such conflicts which mostly occur with a multi-user system. It helps you to make sure that database transactions are performed concurrently without violating the data integrity of respective database. Therefore, Concurrency Control is the most important element for the proper functioning of a multi-user system.

* Concurrency Control protocols.

Different Concurrency control protocols offer different benefits between the amount of concurrency they allow & the amount of overhead that they impose.

- Lock-Based Protocols
- Two Phase
- Timestamp-Based Protocols
- Validation-Based Protocols

Two-Phase Locking Technique

Concept of two-phase locking

Some of the main techniques for concurrency control of transactions are based on the concept of locking data items.

A lock is a mechanism to control concurrent access to a data item. Generally, there is one lock for each data item in the database. Any transaction cannot read or write data until it acquires an appropriate lock on it.

→ Locking is an operation which secures permission to read, or write on a data item. Two phase locking is a process used to gain ownership of shared resources without creating the possibility of deadlock. To guarantee serializability, we use two-phase locking protocol. A transaction is said to follow the Two-Phase Locking protocol if Locking & Unlocking can be done in two phases.

• Growing phase : Here a new lock on the data item may be acquired by the transaction, but none can be released.

• Shrinking Phase : Here, existing lock held by the transaction may be released, but no new locks can be acquired.

The two main modes in which a data item may be locked are

- Exclusive(X) mode \rightarrow Data Item can be read as well as written.
- Shared(S) mode \rightarrow Data Item can only be read.

Request Item	S	X
S	True	False
X	False	False

fig: Lock Compatibility matrix.

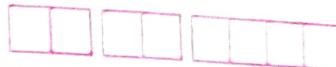
A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transaction.

Types of locks and system locks table

There are various ~~mode~~ kinds of locks

1. Binary Locks

In binary lock, a lock on a data item can be in two states: It is either locked or unlocked. (or 1 and 0, for simplicity). A distinct lock is associated with each database item, X. If the value of the lock on X is 1, item X can't be accessed by a database operation that requests the item. Otherwise, if it is 0, then the item can be accessed when requested. The lock value is changed to 1. We refer to the current state of the lock associated with item X as lock(X). Two operations, lock(item X) & unlock(item X) are used with binary locking.



Q. Shared-Exclusive or (Read/Write) Lock.

This type of locking mechanism differentiates the locks based on their use. If lock is acquired on a data item to perform a write operation, it is an exclusive lock. Allowing more than one transaction to write on the same data item would lead the database into an inconsistent state. Read locks are shared because no data value is being changed.

* Lock Conversion (Upgrading, Downgrading)

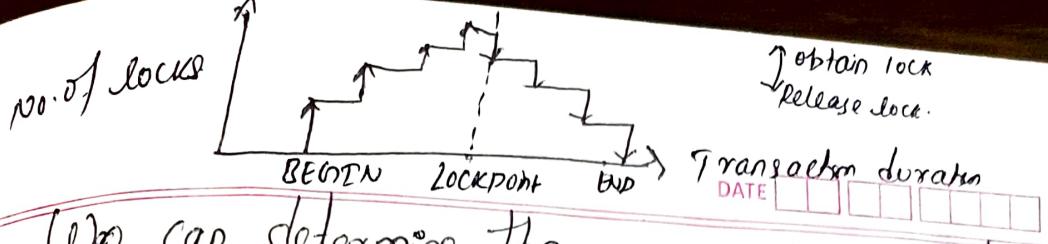
The mechanism of conversion from shared lock to exclusive lock is known as upgrade of lock. Whereas the mechanism of conversion from exclusive lock to shared lock is downgrade of lock. When upgrading and downgrading of locks is used, the lock table must include transaction identifiers in the record structure for each lock to store the information on which the transaction holds locks on the item.

* Disadvantages of two-phase locking or Shared-Exclusive

- May not be free from Precoverability
- May not be free from deadlock
- May not be free from starvation
- May not be free from cascading rollback

* Guaranteeing Serializability by Two-phase locking

A transaction is said to follow the 2P locking protocol if all locking operations (read-lock, write-lock) precede the first unlock operation. Such a transaction can be divided into two phases:-
 Phase 1: during which new locks can be acquired but none can be released; and a shrinking phase: during which existing locks can be released but no new locks can be acquired. If lock conversion is allowed, then upgrading of locks must be done during the expanding phase & downgrading must be done in shrinking phase.



We can determine the serializability on the basis of its lock point. Here lock point ensures that T_1 has occurred first which is followed by T_2 . There is no any loop. Thus it ensures serializability. This is one of the main advantage of 2PL.

	T_1	T_2	
T ₁ growing phase	1 LOCK-S(A)		\Rightarrow Shared lock
	2	LOCK-S(A)	X \Rightarrow Exclusive lock
Lock point of T_1	3 LOCK-X(B)		T ₂ growing phase
	4 -----	-----	
T ₁ shrinking phase	5 UNLOCK(A)		Lock point of T_2
	6	LOCK-X(C)	
T ₂ shrinking phase	7 UNLOCK(B)		T ₂ shrinking phase
	8	UNLOCK(A)	
	9	UNLOCK(C)	
	10 -----	-----	

fig 2PL

Types of two-phase locking

There are four types of two-phase locking

- 1. Basic 2PL \rightarrow we just studied
 - 2. Conservative 2PL
 - 3. Strict 2PL
 - 4. Rigorous 2PL
- } imp.

Conservative 2PL

According to this scheme, a transaction locks all data items before beginning its operations and releases those data items only after executing its last operation i.e. pre-declare all the locks on all variables before its execution. It solves the problem of deadlock, irrecoverability, starvation and cascading rollback but it often degrades the system performance by unnecessary delaying other transactions which want to access some of the data item locked by the transaction. It is very difficult to use in practice.

Strict 2PL

According to this protocol, a transaction does not release the exclusive locks i.e. write lock until it commits or aborts. It solves the irrecoverability problem & problem of cascading rollback but not the problem of deadlock & starvation.

Note: It should satisfy the basic 2PL

T ₁	T ₂	T ₃
----------------	----------------	----------------

X(A)

R(A)

W(A)

Synclock(A)

-

-

S(A)

R(A)

S(A)

R(A)

X(A)

R(A)

W(A)

T₂

-

-

>S(A)

R(A)

T₃

-

-

S(A)

R(A)

rollback

classmate

Fair

cascading rollback

problem

Commit
unlock(A)PAGE

--	--	--

fig: solving rollback problem

Note:- Cascading rollback problem :- If the failure of one transaction causes several other dependent transactions to rollback or abort, pt is called crp.

Irrecoverability :- The failure of a trans action has cause another transaction to rollback. But the another transaction has already committed, so pt is Prerecoverable.

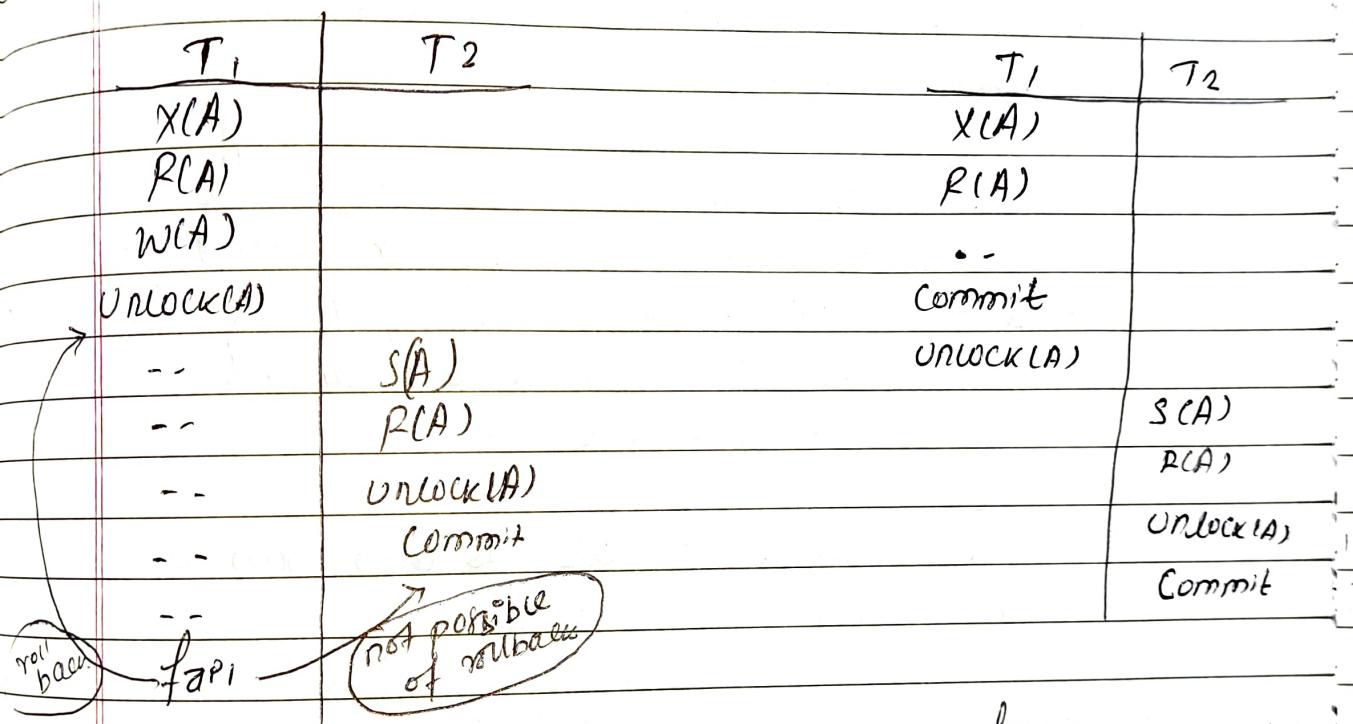


fig: Problem of prerecoverability

fig: Irrecoverable problem solved

Rigorous 2PL

According to this protocol, a transaction does not release any of exclusive as well as shared locks until it commits or aborts. It also solves the problem of Cascading rollback & irrecoverability but the problem of deadlock & starvation still exists.

Dealing with deadlock & starvation

Deadlock

1) A deadlock is a condition in which two or more transaction are waiting for each other.

2) It occurs when two or more transactions are unable to proceed because each is waiting for one of the other to do some thing.

3) It is also known as Circular waiting.

Avoidance:

- Acquire locks at once before starting
- Acquire locks with predefined order.

6) Resources are blocked by the processes

Starvation

1) Starvation happens if same transaction is always chosen as victim.

2) It occurs if the waiting scheme for locked items is unfair, giving priority to some transactions over others.

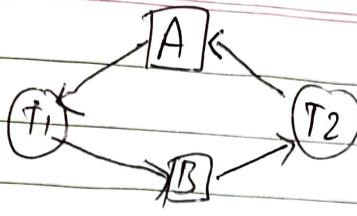
3) It is also known as lived lock.

Avoidance:

- Switch priorities so that every thread has a chance to have high priority
- Use FIFO order among competing requests.

7) Resources are continuously utilized by high priority processes

Deadlock



Two transaction T_1 & T_2 need two Resource A & B for completing the task. T_1 requests resource A & is granted and holds it. T_2 requests resource B & is also granted & holds it. Now, if T_1 request for B held by T_2 , then request is denied until T_2 release it. Unfortunately, instead of releasing B, T_2 asks for A held by T_1 . At this state, both transaction are blocked & will remain so forever. This situation is an example of deadlock.

* Necessary Conditions for deadlock.

- Mutual exclusion \rightarrow Only one transaction at a time can use a resource
- Hold & Wait \rightarrow Transaction holding a resource while waiting to acquire another resource held by another transaction
- No preemption \rightarrow Resources are only released when a transaction is executed
- Circular wait \rightarrow One transaction T_i waiting for resource held by another transaction T_j while T_j is waiting for resource held by T_k & T_k is waiting for resource held by T_i .

* Deadlock detection

Maintain a wait-for graph i.e. resource allocation graph where node represent processes and if process P_i is waiting for a resource held by process P_j , then there exists an edge from i to j . If there is a cycle in the graph, deadlock exists.



1 Deadlock Recovery

When a deadlock algorithm determines that a deadlock exists, the system must recover from deadlock. There are two approaches:

- Simply kill one or more processes to break circular wait.
- Preempt some resources from one or more processes involved in deadlock.

② Dealing with Starvation

One solution for starvation is to have a fair waiting scheme, such as using a first-come-first-served queue, where transactions are enabled to lock an item in order in which they originally requested to lock. Another scheme allows some transactions to have higher priority over others. The wait-die and wound-wait avoid starvation, because they restart a transaction that have been aborted.

(Timestamp Ordering)

Timestamp ordering Concurrency Control concept

A timestamp is a unique identifier created by the DBMS to identify a transaction. The timestamp ordering is used to order the transactions based on their Timestamps. The order of transaction is nothing but the ascending order of the transaction creation. The priority of the older transaction is high so it executes first. To determine the timestamp of a transaction, the system uses ~~the time is~~ by the protocol.

Let's assume there are two transactions T_1 and T_2 . Suppose the transaction T_1 has entered the system at 007 times & T_2 has entered at 009. T_1 has the higher priority, so it executes 1st as it entered the system 1st. The timestamp ordering protocol also maintains the timestamp of last 'read' and 'write' operation on data.

Timestamp values are assigned in order in which the transactions are submitted to the system, so it can be thought of as the transaction start time. We denote timestamp transaction of transaction T as $TS(T)$. Concurrency control techniques based on timestamp ordering don't use locks, hence, deadlocks can't occur.

- Read timestamp (R-TS) :- Latest transaction number which performed read successfully
- Write timestamp (W-TS) :- Latest transaction number which performed write successfully

* Basic Time Stamp ordering:

Rules:-

1) Transaction T_i issues a Read(A) operation

- a) If $WTS(A) > TS(T_i)$, Rollback T_i (operation rejected)
- b) Otherwise execute $R(A)$ operation
Set $RTS(A) = \max\{RTS(A), TS(T_i)\}$

2) Transaction T_i issues Write(A) operation

- a) If $RTS(A) > TS(T_i)$ then Rollback T_i (operation rejected)
- b) If $WTS(A) > TS(T_i)$ then Rollback T_i (" ")
- c) Otherwise execute Write(A) operation
Set $WTS(A) = TS(T_i)$

* Strict Time stamp ordering

A variation of basic Time stamp ordering called Strict $T_0^{(Time)}$ ensure that the schedules are both strict (for easy recoverability) & (conflict) serializable

1. Transaction T_i issues a write-item (A) operation:

- If $TS(T) > RTS(A)$, then delay T until the transaction T' that wrote or read A has committed or aborted

2. (Read-item)_(A) • If $TS(T) > WTS(A)$, then delay T until the transaction T' that wrote or read A has committed or aborted.

Thomas's Write rule

A modification of the basic TO algorithm, known as Thomas's write rule, does not enforce conflict serializability, but it rejects fewer write operations by modifying the checks for the write-item (A) operation as follows.

- If $PTS(A) > TS(T)$, then abort & rollback T & reject the operation
- If $WTS(A) > TS(T)$, then just ignore the write operation and continue execution. This is because the most recent writer counts in case of two consecutive writes
- If the conditions given in above two do not occur, then execute write-item(A) of T & set $WTS(A)$ to $TS(T)$

(Multiversion Concurrency Control)

* Concept of multiversion concurrency control technique.

- It provides concurrent access to database without locking the data.
- This technique keeps the record of old values of a data item when the data item is updated.
- These are known as multiversion concurrency control, because several version (values) of an item are maintained.
- When a transaction requires access to an item an appropriate version is chosen to maintain the serializability of the current executing schedule, if possible.
- It improves the performance of database applications in a multiuser environment.

* Multiversion technique based on timestamp ordering

Assume x_1, x_2, \dots, x_n are the versions of a data item x created by a write operation of transaction. For each version, the value of version x_i , a RTS & WTS are associated.

- RTS: The read timestamp r_s is the largest of all the timestamps of transactions that have successfully Read Version.
- WTS: The write timestamp is the timestamp of the transaction that wrote the value of version.

To ensure Serializability, the following two rules are as:

- If transaction T issues a write item(x) operation, & version of x has highest WTS() of all versions of x that is also less than or equal to TS(T), & PTS() > TS(T), then abort & roll back transaction T; otherwise create a new version of x with RTS() = WTS() = TS(T)
- If transaction T issues a Read(x) operation, find version of x that has highest WTS() of all versions of x that is also less than or equal to TS(T); then return the value of to transaction T, & set the value of RTS() to the larger of TS(T) & the current PTS()

Multiversion locking using certify locks (2 Phase locking)

In this scheme, there are three locking modes for an item; Read, Write & Certify. Hence, the state of Lock(X) for an item X can be one of read-locked, write-locked, certify-locked or unlocked. The idea behind it is to allow other transactions T_2 to read an item X while a single transaction T_1 holds a write lock on X. This is accomplished by allowing two versions for each item X; one, the committed version, must always have been written by some committed transaction. and the second local version X' can be created when a transaction T acquires a write lock on X.

Once T is ready to commit, it must obtain a certify lock on all items that it currently holds

write locks on before Pt can commit. The certify lock is not compatible with read locks, so the transaction may have to delay its commit until all Pts write-locked items are released by any reading transactions in order to obtain the certify locks. Once the certify locks are acquired, the committed version x of the data item is set to the value of version x , version x is discarded, & the certify locks are then released.

(Validation (optimistic) Techniques & Snapshot Isolation (Concurrency Control))

* Concept of validation (optimistic) Techniques

The locked based & time stamp ordering protocols were pessimistic techniques. Because they always think negative that the transactions can go in conflict.

Validation phase is a optimistic control technique. Here, the transaction is executed in 3 phases

- Read Phase: Here, the transaction T is read & executed. Pt is used to read the value of various data item & stores them in temporary local variables. T can perform all the write operations on temporary variables without an update to the actual database.

- Validation phase:- Here, temporary variable value will be validated against the actual data to see if it violates the serializability.
- Write phase:- If the validation of the transaction is validated, then the temporary results are written to the database or system otherwise the transaction is rolled back.
(ie. the changes are permanently applied to the DB)

Note: Each phase has following timestamps.

→ Start (T_i^1):- time when T_i started execution

→ Validation (T_i^2):- time when T_i finishes its read phase & starts its validation phase.

→ Finish (T_i^3):- time when T_i finishes its write phase

* Snapshot Isolation Concurrency Control.

- Here, each transaction operates on database Snapshot.
- The Snapshot is taken once the transaction starts.
- When transaction concludes, it will successfully commit only if the values update by the transaction haven't been changed externally since Snapshot was taken.
- In practice, Snapshot Isolation is implemented within multi version concurrency control (Mvcc), where generational values of each data item (versions) are maintained.
- It is widely used in practice.
- The main reason for its adoption is that it allows better performance than serializability, yet still avoids most of the concurrency anomalies.

Serializability avoids (but not always all).

Questions asked from this chapter



- Q. Explain deadlock and starvation. Explain Time stamp based protocol for concurrency control. (2078-10 marks) (2075-5 marks)
- Q. What is Concurrency control? Name various methods of controlling the Concurrency control. Differentiate between Binary lock and Shared/Exclusive lock (2076-10 marks) (2073-6 marks)
- Q. What is Granularity of data items? How does it effect in Concurrency control? (2069-5 marks) (2067-5 marks) (2067-5 marks)
- Q. Explain 2 phase locking technique in brief. (2076-5 marks) (2072-6 marks) (2067-5 marks) (2068-5 marks)
- Q. Discuss the problems of deadlock & starvation, and the different approaches to dealing with these problems. (2067-5 marks) (2073-5 marks) (2025-6 marks)